

## Extended Abstract Track

# Weight-space gradient descent induces approximate activity descent, making learning rules testable

**Editors:** List of editors' names

## Abstract

A central question in neuroscience is whether biological neural networks optimise a loss function by adjusting synaptic weights – and if so, by what rule. Decades of theoretical work have proposed candidate learning rules. But synaptic weights cannot be measured at scale in behaving animals, so these proposals remain difficult to test directly. Neural *activities*, by contrast, can be recorded in large populations across learning. This raises a concrete question: if the brain does perform gradient descent on its weights, what should we expect to see in the activities? Here we show that the answer takes a surprisingly structured form: weight-space gradient descent induces kernel-filtered descent in activity space, and in diagonal-dominant regimes this reduces to approximately per-neuron activity descent. The resulting prediction is directly testable with existing experimental techniques.

**Keywords:** neural activity, gradient descent, neural tangent kernel, learning rules, computational neuroscience

## 1. Weight-space descent and observable activity

Gradient descent on a neural network embodies a simple learning principle: *the more that strengthening a synapse would improve downstream performance, the more that synapse is strengthened* (and vice versa—harmful synapses are weakened). We consider what this principle implies for the *observable* neural activities in a network that is performing weight-space gradient descent.

**Setup and derivation.** Consider a feedforward network with layers  $\ell = 1, \dots, L$ , weights  $\mathbf{W}$ , and neural activities  $A_i^{(\ell)}(\mathbf{W})$ , trained to minimise a loss  $\mathcal{L}$ . Let  $\mathcal{N}_\ell := \{1, \dots, n_\ell\}$  index the neurons in layer  $\ell$ , let  $\mathcal{P}_r$  index the scalar weights in layer  $r$ , and define the upstream parameter-index set  $\mathcal{P}_{\leq \ell} := \bigcup_{r=1}^{\ell} \mathcal{P}_r$ . Thus  $i, j, k \in \mathcal{N}_\ell$  index neurons in a fixed layer, whereas  $\alpha \in \mathcal{P}_{\leq \ell}$  indexes a scalar upstream weight  $W_\alpha$ . We write  $\mathbf{W}_{\leq \ell} := (W_\alpha)_{\alpha \in \mathcal{P}_{\leq \ell}}$  for the corresponding block of weights.

In gradient descent, every weight is updated simultaneously:

$$\Delta W_\alpha = -\eta \frac{\partial \mathcal{L}}{\partial W_\alpha}, \quad \forall \alpha \in \mathcal{P}_{\leq L}. \quad (1)$$

What is the induced change in the activity of neuron  $i$  in layer  $\ell$ ? Because  $A_i^{(\ell)}$  depends only on  $\mathbf{W}_{\leq \ell}$ , if the activity map is twice differentiable with bounded second derivatives along the update segment, Taylor expansion around the pre-update weights  $\mathbf{W}_t$  gives

$$\Delta A_i^{(\ell)} := A_i^{(\ell)}(\mathbf{W}_t + \Delta \mathbf{W}) - A_i^{(\ell)}(\mathbf{W}_t) = \nabla_{\mathbf{W}_{\leq \ell}} A_i^{(\ell)}(\mathbf{W}_t)^\top \Delta \mathbf{W}_{\leq \ell} + O(\|\Delta \mathbf{W}_{\leq \ell}\|^2).$$

Since  $\Delta \mathbf{W} = O(\eta)$ , substituting  $\Delta W_\alpha = -\eta \partial \mathcal{L} / \partial W_\alpha$  gives

$$\Delta A_i^{(\ell)} \approx -\eta \sum_{\alpha \in \mathcal{P}_{\leq \ell}} \frac{\partial A_i^{(\ell)}}{\partial W_\alpha} \frac{\partial \mathcal{L}}{\partial W_\alpha}, \quad (2)$$

# Extended Abstract Track

where all derivatives are evaluated at  $\mathbf{W}_t$  and the omitted remainder is  $O(\eta^2)$ . For piecewise-linear activations such as ReLU, this statement applies when the step is small enough that no unit changes activation status, so the network remains within a single smooth piecewise-polynomial region of weight space. In a sequential feedforward network (one without skip or residual connections), the effect of any weight indexed by  $\mathcal{P}_{\leq \ell}$  on the loss passes entirely through the activities of layer  $\ell$ . This allows us to use the chain rule to write the effect on the loss,  $\partial \mathcal{L} / \partial W_\alpha$ , in terms of the activities at layer  $\ell$ :

$$\Delta A_i^{(\ell)} \approx -\eta \sum_{\alpha \in \mathcal{P}_{\leq \ell}} \frac{\partial A_i^{(\ell)}}{\partial W_\alpha} \left[ \sum_{k \in \mathcal{N}_\ell} (\partial \mathcal{L} / \partial A_k^{(\ell)}) (\partial A_k^{(\ell)} / \partial W_\alpha) \right], \quad (3)$$

Exchanging the sums yields

$$\Delta A_i^{(\ell)} \approx -\eta \sum_{k \in \mathcal{N}_\ell} \Phi_{ik}^{(\ell)} \frac{\partial \mathcal{L}}{\partial A_k^{(\ell)}}, \quad (4)$$

where

$$\Phi_{ik}^{(\ell)} := \sum_{\alpha \in \mathcal{P}_{\leq \ell}} \frac{\partial A_i^{(\ell)}}{\partial W_\alpha} \frac{\partial A_k^{(\ell)}}{\partial W_\alpha} = \nabla_{\mathbf{W}_{\leq \ell}} A_i^{(\ell)} \cdot \nabla_{\mathbf{W}_{\leq \ell}} A_k^{(\ell)} \quad (5)$$

is a Gram matrix defined on the activities within each layer, and  $\partial \mathcal{L} / \partial A_k^{(\ell)}$  is the full backpropagated loss gradient – a sufficient statistic that compresses everything downstream layers contribute to the weight updates affecting layer  $\ell$ . In vector form,  $\Delta \mathbf{A}^{(\ell)} \approx -\eta \Phi^{(\ell)} \nabla_{\mathbf{A}^{(\ell)}} \mathcal{L}$ : gradient descent in weight space induces *kernel descent* within each layer of activity space. It holds for any differentiable sequential feedforward network regardless of depth, width, or activation function. Architectures with skip connections require additional treatment, as weights can influence the loss through paths that bypass intermediate layers.

**Relation to the neural tangent kernel.** The first-order approximation we study should be familiar to those acquainted with the neural tangent kernel (NTK) (Jacot et al., 2018). The kernel  $\Phi^{(\ell)}$  is a layer-wise analogue of the empirical neural tangent kernel (NTK) for finite-width networks (Lee et al., 2020). We do not assume the network is wide enough for our kernel to be deterministic, and it need not be so for our purposes.

**The self term drives each neuron toward its own loss gradient.** The diagonal element  $\Phi_{ii}^{(\ell)} = \|\nabla_{\mathbf{W}_{\leq \ell}} A_i^{(\ell)}\|^2$  collects contributions from *all* weights that influence neuron  $i$ —both its own incoming weights and every weight in earlier layers that lies on a path to it. Because  $\Phi_{ii}^{(\ell)} \geq 0$ , the self-term update is aligned with  $-\partial \mathcal{L} / \partial A_i^{(\ell)}$ : each such neuron’s activity moves down its own loss gradient, with an effective learning rate  $\eta \Phi_{ii}^{(\ell)}$  that grows with the depth and width of its input pathways. Dead ReLU units have  $\Phi_{ii}^{(\ell)} = 0$  and contribute no self-term drive.

**Cross terms introduce interference from other neurons.** The off-diagonal elements  $\Phi_{ik}^{(\ell)} = \nabla_{\mathbf{W}_{\leq \ell}} A_i^{(\ell)} \cdot \nabla_{\mathbf{W}_{\leq \ell}} A_k^{(\ell)}$  for  $k \neq i$  couple neuron  $i$  to every other neuron’s loss gradient *within the same layer*. Their magnitude reflects similarity of weight-space sensitivities,

## Extended Abstract Track

often arising from shared upstream parameters: two neurons whose activities depend on overlapping upstream weight directions can have large  $|\Phi_{ik}^{(\ell)}|$ , whereas neurons supported by effectively orthogonal parameter directions have  $\Phi_{ik}^{(\ell)} \approx 0$ . These cross terms have no reason to be aligned with  $-\partial\mathcal{L}/\partial A_i^{(\ell)}$  and constitute interference—the side-effect of the rest of the layer learning simultaneously. Whether the self term or the cross terms dominate determines whether activity changes look like activity-space gradient descent.

**Approximate kernel diagonality.** The cross terms are small when different neurons’ activities depend on nearly orthogonal directions in upstream weight space. Equivalently,

$$\nabla_{\mathbf{W}_{\leq \ell}} A_i^{(\ell)} \cdot \nabla_{\mathbf{W}_{\leq \ell}} A_k^{(\ell)} \approx 0, \quad i \neq k.$$

Whether this occurs depends on the architecture, the parameterization, and the learning task.

*Width.* Increasing width gives each neuron more private incoming parameters, which can enlarge the self term  $\Phi_{ii}^{(\ell)}$ . But neurons in the same layer also share all upstream parameters, so their activity gradients are generally not independent. In some architectures and training regimes, width may decorrelate the shared component of the gradients, but this is not universal.

*Depth.* Deeper layers share more upstream parameters. This can strengthen correlations between neurons’ activity gradients.

*Learning dynamics.* Training may also change this balance. If learning decorrelates representations or routes different neurons through increasingly distinct effective pathways, the off-diagonal entries of  $\Phi^{(\ell)}$  may shrink relative to the diagonal.

**The diagonal approximation: activity descent as a proxy for weight descent.** Combining the kernel identity (Eq. 3) with diagonal dominance yields a simple approximation:

$$\Delta A_i^{(\ell)} \approx -\eta \Phi_{ii}^{(\ell)} \frac{\partial \mathcal{L}}{\partial A_i^{(\ell)}}. \quad (6)$$

The intuitive meaning mirrors that of weight descent: just as weight descent says *strengthen each synapse in proportion to how much it helps*, activity descent says *increase each neuron’s firing in proportion to how much that increase would help* (and decrease it when it hurts). Equation 5 therefore identifies a useful approximation regime: when  $\Phi^{(\ell)}$  is sufficiently diagonal-dominant, each neuron’s activity moves in the direction that locally reduces the loss, up to a neuron-specific scaling factor  $\Phi_{ii}^{(\ell)}$ .

**A testable prediction for neuroscience.** This result converts an untestable claim about synaptic learning rules into a testable prediction about observable activity changes. Synaptic weights cannot be measured at scale in vivo, but neural activities can. If a biological circuit performs *any* form of gradient descent on its synaptic weights—even an approximate or noisy variant—and operates in a regime where  $\Phi^{(\ell)}$  is approximately diagonal-dominant, then the observable activity changes must satisfy Eq. 5. The per-neuron scaling  $\Phi_{ii}^{(\ell)}$  is unknown and may vary across neurons, but the *sign* of each active neuron’s activity change is pinned: because  $\Phi_{ii}^{(\ell)} \geq 0$  (and is strictly positive for active units), every such neuron must move in the direction that locally reduces the loss, i.e.

# Extended Abstract Track

$\text{sign}(\Delta A_i^{(\ell)}) \approx \text{sign}(-\partial \mathcal{L} / \partial A_i^{(\ell)})$ . Note that if  $\Phi_{ii}^{(\ell)}$  varies substantially across neurons, the population vector  $\Delta \mathbf{A}^{(\ell)}$  will not be parallel to  $-\nabla_{\mathbf{A}^{(\ell)}} \mathcal{L}$  but rather to  $-\text{diag}(\Phi_{ii}^{(\ell)}) \nabla_{\mathbf{A}^{(\ell)}} \mathcal{L}$ ; the two directions coincide only when  $\Phi_{ii}^{(\ell)}$  is approximately uniform. A concrete experimental protocol follows: record neural population activity across learning trials, estimate  $\partial \mathcal{L} / \partial A_i^{(\ell)}$  from a task-fitted model (?), and check whether the measured  $\Delta A_i^{(\ell)}$  covaries positively with  $-\partial \mathcal{L} / \partial A_i^{(\ell)}$  across neurons. Experimentally,  $\Delta A_i^{(\ell)}$  should be estimated for matched task conditions or repeated stimuli across learning, so that the measured change is attributable to learning rather than to changes in the input distribution or behavioral state. The unknown  $\Phi_{ii}^{(\ell)}$  absorbs per-neuron scaling, so the test checks the sign pattern and rank-order correlation, not the exact magnitudes. A robust failure of this sign alignment would falsify the conjunction of assumptions that the fitted downstream model gives the correct activity-space loss gradient, that the circuit is following a gradient-like synaptic update for that loss, and that the relevant activity kernel is sufficiently diagonal-dominant. Because this test relies on estimating  $\partial \mathcal{L} / \partial A_i^{(\ell)}$  from a fitted downstream model, a negative result is ambiguous unless that model can be independently validated.

## **Caveat: activity descent is necessary but not sufficient for weight descent.**

The converse is weaker: observing that activities follow their loss gradient constrains the projection of  $\Delta \mathbf{W}$  onto the row space of the activity Jacobian  $J_A = \partial \mathbf{A} / \partial \mathbf{W}$ , but says nothing about the (typically vast) null space of weight changes that leave activities unchanged on the current input. Many weight updates can induce the same observed  $\Delta \mathbf{A}$ ; the minimum-norm update for a specified activity change is given by the Jacobian pseudoinverse,  $\Delta \mathbf{W}_{\min} = J_A^\top (J_A J_A^\top)^+ \Delta \mathbf{A}$ , not generally by backpropagation. Vanilla weight-gradient descent  $\Delta \mathbf{W}_{\text{GD}} = -\eta J_A^\top \nabla_{\mathbf{A}} \mathcal{L}$  is one particular row-space update: it induces kernel descent  $\Delta \mathbf{A} \approx -\eta J_A J_A^\top \nabla_{\mathbf{A}} \mathcal{L}$ , but the activity change alone does not identify that weight-level rule. In short, weight descent implies the first-order kernel relation in Eq. 3, and implies the simpler activity-descent form in Eq. 5 when  $\Phi^{(\ell)}$  is close to diagonal.

## **2. Experiments.**

Figures 1–2 numerically verify Eqs. 2–5 in a 3-hidden-layer ReLU MLP trained with online SGD ( $\eta = 0.005$ ) on a 2-class image dataset (4×4 pixel images of horizontal vs. vertical bars, 20 per class, input dimension 16, output dimension 2). Dead ReLU neurons (pre-activation  $\leq 0$  on the sampled input) are excluded from all analyses. At regular intervals during training, we compare predicted activity changes against the actual change measured after a single gradient step. Figure 1 computes the full activity Jacobian  $\partial A_i^{(\ell)} / \partial W_\alpha$  and the per-layer kernel  $\Phi_{ik}^{(\ell)}$  to show detailed diagnostics at widths 8 and 48. Figure 2 computes the same kernel diagonal exactly via the layerwise kernel recursion and sweeps width from 4 to 512 (3 random seeds per width, each trained for 2000 SGD steps) to quantify how the diagonal approximation in Eq. 5 improves with width in this toy setting.

## Extended Abstract Track




Figure 1: **In a wide network doing weight descent, activity changes closely match activity-space gradient descent.** Top row: width=8. Middle row: width=48. Left: unnormalised  $\Phi_{ik}^{(\ell)}$  heatmaps, each with its own horizontal colorbar (white = 0; colorbar shows raw kernel magnitude). Scatter panels: predicted vs. actual  $\Delta A$  for the full kernel (Eq. 3), the diagonal approximation  $-\Phi_{ii} \partial \mathcal{L} / \partial A_i$ , and raw  $-\partial \mathcal{L} / \partial A$ . Dead ReLU neurons excluded. Bottom row:  $r(\Delta A, -\Phi_{ii} \partial \mathcal{L} / \partial A_i)$  and loss over training.

# Extended Abstract Track




Figure 2: **The diagonal approximation improves with width in this toy setting.** Pearson  $r$  between actual  $\Delta A$  and  $-\Phi_{ii} \partial \mathcal{L} / \partial A_i$  as a function of hidden-layer width ( $\log_2$  scale). Dead ReLU neurons excluded. Shaded regions:  $\pm 1$  s.d. across seeds.

## References

- Arthur Jacot, Franck Gabriel, and Clément Hongler. Neural Tangent Kernel: Convergence and Generalization in Neural Networks, June 2018.
- Jaehoon Lee, Lechao Xiao, Samuel S. Schoenholz, Yasaman Bahri, Roman Novak, Jascha Sohl-Dickstein, and Jeffrey Pennington. Wide Neural Networks of Any Depth Evolve as Linear Models Under Gradient Descent. *Journal of Statistical Mechanics: Theory and Experiment*, 2020(12):124002, December 2020. ISSN 1742-5468. doi: 10.1088/1742-5468/abc62b.

## Extended Abstract Track

**Appendix A. Annotated layer-2 decomposition**

To make the diagonal/off-diagonal split concrete, consider a standard layer-2 block (omitting biases for simplicity):

$$z_i^{(2)} = \sum_m W_{im}^{(2)} A_m^{(1)}, \quad A_i^{(2)} = \sigma_2(z_i^{(2)}).$$

The first-order activity update is

$$\Delta A_i^{(2)} \approx -\eta \sum_j \Phi_{ij}^{(2)} \frac{\partial \mathcal{L}}{\partial A_j^{(2)}}.$$

The point of this appendix is to show explicitly where the self term  $\Phi_{ii}^{(2)}$  and the cross terms  $\Phi_{ij}^{(2)}$  for  $j \neq i$  come from.

**Own-layer weights contribute only to the diagonal.** For a weight  $W_{pm}^{(2)}$  feeding directly into layer 2,

$$\frac{\partial A_i^{(2)}}{\partial W_{pm}^{(2)}} = \delta_{ip} \sigma_2'(z_i^{(2)}) A_m^{(1)}.$$

Therefore the layer-2 contribution to the kernel is

$$\sum_{p,m} \frac{\partial A_i^{(2)}}{\partial W_{pm}^{(2)}} \frac{\partial A_j^{(2)}}{\partial W_{pm}^{(2)}} = \delta_{ij} [\sigma_2'(z_i^{(2)})]^2 \|\mathbf{A}^{(1)}\|^2 = \delta_{ij} D_{2,ii}.$$

This term is diagonal because two different neurons in the same layer do not share their own incoming weights.

**Earlier-layer weights are propagated forward by the Jacobian.** Now take a weight  $W_\alpha$  with  $\alpha \in \mathcal{P}_1$ . By the chain rule,

$$\frac{\partial A_i^{(2)}}{\partial W_\alpha} = \sum_m \frac{\partial A_i^{(2)}}{\partial A_m^{(1)}} \frac{\partial A_m^{(1)}}{\partial W_\alpha} = \sum_m J_{2,im} \frac{\partial A_m^{(1)}}{\partial W_\alpha}, \quad J_{2,im} = \frac{\partial A_i^{(2)}}{\partial A_m^{(1)}}.$$

Summing over all layer-1 weights gives

$$\sum_{\alpha \in \mathcal{P}_1} \frac{\partial A_i^{(2)}}{\partial W_\alpha} \frac{\partial A_j^{(2)}}{\partial W_\alpha} = \sum_{m,n} J_{2,im} \Phi_{mn}^{(1)} J_{2,jn}.$$

Because layer 1 has no earlier shared weights,  $\Phi^{(1)} = D_1$  is diagonal, so this reduces to

$$\sum_{m,n} J_{2,im} \Phi_{mn}^{(1)} J_{2,jn} = \sum_m J_{2,im} D_{1,mm} J_{2,jm}.$$

Thus the full layer-2 kernel is

$$\Phi_{ij}^{(2)} = \delta_{ij} D_{2,ii} + \sum_m J_{2,im} D_{1,mm} J_{2,jm}.$$

# Extended Abstract Track

If one writes the Jacobian explicitly,

$$J_{2,im} = \sigma'_2(z_i^{(2)}) W_{im}^{(2)},$$

so

$$\Phi_{ij}^{(2)} = \delta_{ij} [\sigma'_2(z_i^{(2)})]^2 \|\mathbf{A}^{(1)}\|^2 + \sigma'_2(z_i^{(2)}) \sigma'_2(z_j^{(2)}) \sum_m W_{im}^{(2)} D_{1,mm} W_{jm}^{(2)}.$$

**Self and cross terms in the activity update.** Substituting the layer-2 kernel into the first-order update and separating the  $j = i$  and  $j \neq i$  pieces gives

$$\Delta A_i^{(2)} \approx -\eta \Phi_{ii}^{(2)} \frac{\partial \mathcal{L}}{\partial A_i^{(2)}} - \eta \sum_{j \neq i} \Phi_{ij}^{(2)} \frac{\partial \mathcal{L}}{\partial A_j^{(2)}}.$$

Using the expression above,

$$\Phi_{ii}^{(2)} = D_{2,ii} + \sum_m J_{2,im}^2 D_{1,mm},$$

while for  $j \neq i$ ,

$$\Phi_{ij}^{(2)} = \sum_m J_{2,im} D_{1,mm} J_{2,jm}.$$

Hence

$$\Delta A_i^{(2)} \approx -\eta \left( D_{2,ii} + \sum_m J_{2,im}^2 D_{1,mm} \right) \frac{\partial \mathcal{L}}{\partial A_i^{(2)}} - \eta \sum_{j \neq i} \left( \sum_m J_{2,im} D_{1,mm} J_{2,jm} \right) \frac{\partial \mathcal{L}}{\partial A_j^{(2)}}.$$

The first bracket is the self term: it is always aligned with  $-\partial \mathcal{L} / \partial A_i^{(2)}$ . The second sum is the interference term: it mixes neuron  $i$  with the gradients of other neurons in the same layer through shared earlier-layer parameters.

**Relation to the temporary  $x_0/x_1$  notation.** If one writes  $x_0 = \mathbf{A}^{(0)}$  and  $x_1 = \mathbf{A}^{(1)}$ , then

$$D_{1,mm} = [\sigma'_1(z_m^{(1)})]^2 \|x_0\|^2, \quad D_{2,ii} = [\sigma'_2(z_i^{(2)})]^2 \|x_1\|^2.$$

So the appearance of both  $x_0$  and  $x_1$  in scratch work is meaningful, not a typo: they are simply the activities entering layers 1 and 2. In the main text,  $\mathbf{A}^{(0)}$  and  $\mathbf{A}^{(1)}$  are cleaner and more consistent with the rest of the notation.

## Appendix B. General recursion for deeper networks and nonlinearities

For an arbitrary depth- $L$  feedforward network with pointwise nonlinearities,

$$z_i^{(\ell)} = \sum_m W_{im}^{(\ell)} A_m^{(\ell-1)}, \quad A_i^{(\ell)} = \sigma_\ell(z_i^{(\ell)}), \quad \ell = 1, \dots, L,$$

again omitting biases for simplicity. The first-order kernel relation from the main text still holds:

$$\Delta \mathbf{A}^{(\ell)} \approx -\eta \Phi^{(\ell)} \nabla_{\mathbf{A}^{(\ell)}} \mathcal{L}, \quad \Phi_{ij}^{(\ell)} = \sum_{\alpha \in \mathcal{P}_{\leq \ell}} \frac{\partial A_i^{(\ell)}}{\partial W_\alpha} \frac{\partial A_j^{(\ell)}}{\partial W_\alpha}.$$

To obtain a recursion, split the weight sum into layer  $\ell$  itself and layers  $1, \dots, \ell - 1$ .



## Extended Abstract Track

**Own-layer contribution.** For a weight  $W_{pm}^{(\ell)}$  in layer  $\ell$ ,

$$\frac{\partial A_i^{(\ell)}}{\partial W_{pm}^{(\ell)}} = \delta_{ip} \sigma'_\ell(z_i^{(\ell)}) A_m^{(\ell-1)}.$$

Therefore the own-layer contribution is diagonal:

$$\sum_{p,m} \frac{\partial A_i^{(\ell)}}{\partial W_{pm}^{(\ell)}} \frac{\partial A_j^{(\ell)}}{\partial W_{pm}^{(\ell)}} = \delta_{ij} [\sigma'_\ell(z_i^{(\ell)})]^2 \|\mathbf{A}^{(\ell-1)}\|^2 = \delta_{ij} D_{\ell,ii}.$$

**Earlier-layer contribution.** Define the inter-layer Jacobian

$$J_{\ell,im} = \frac{\partial A_i^{(\ell)}}{\partial A_m^{(\ell-1)}} = \sigma'_\ell(z_i^{(\ell)}) W_{im}^{(\ell)}.$$

Then for any weight  $W_\alpha$  with  $\alpha \in \mathcal{P}_{\leq \ell-1}$ ,

$$\frac{\partial A_i^{(\ell)}}{\partial W_\alpha} = \sum_m J_{\ell,im} \frac{\partial A_m^{(\ell-1)}}{\partial W_\alpha}.$$

Summing over earlier-layer weights gives

$$\sum_{\alpha \in \mathcal{P}_{\leq \ell-1}} \frac{\partial A_i^{(\ell)}}{\partial W_\alpha} \frac{\partial A_j^{(\ell)}}{\partial W_\alpha} = \sum_{m,n} J_{\ell,im} \Phi_{mn}^{(\ell-1)} J_{\ell,jn}.$$

Combining both pieces yields the recursion

$$\Phi^{(\ell)} = D_\ell + J_\ell \Phi^{(\ell-1)} J_\ell^T, \quad \Phi^{(1)} = D_1.$$

This is the desired recursion in expanded form.

**Unrolling the recursion.** Repeated substitution shows that every earlier layer contributes a positive-semidefinite term transported forward by the intermediate Jacobians:

$$\Phi^{(\ell)} = D_\ell + J_\ell D_{\ell-1} J_\ell^T + J_\ell J_{\ell-1} D_{\ell-2} J_{\ell-1}^T J_\ell^T + \cdots + J_\ell J_{\ell-1} \cdots J_2 D_1 J_2^T \cdots J_{\ell-1}^T J_\ell^T.$$

This makes the structure transparent:  $D_r$  captures the direct effect of layer- $r$  weights on layer- $r$  activities, and the surrounding Jacobians propagate that contribution forward to layer  $\ell$ .

**Role of nonlinearities.** Nothing in the kernel definition or recursion requires linear hidden units. For smooth nonlinearities such as sigmoid, tanh, or GELU, the algebraic derivation above holds with the corresponding  $\sigma'_\ell$ ; the overall  $\Delta \mathbf{A} \approx -\eta \Phi \nabla_{\mathbf{A}} \mathcal{L}$  remains a first-order approximation in  $\eta$ . For piecewise-linear nonlinearities such as ReLU or leaky ReLU, the same formulas hold almost everywhere; at kink points one may use a subgradient. In particular, if a ReLU unit is inactive on the sampled input, then  $\sigma'_\ell(z_i^{(\ell)}) = 0$ , so its row of  $J_\ell$  and its diagonal entry in  $D_\ell$  vanish. This is the appendix-level version of why dead ReLU units contribute nothing to the kernel on that input.

# Extended Abstract Track

**More general feedforward modules.** The recursion above assumes an affine layer followed by a pointwise nonlinearity. The kernel definition

$$\Phi_{ij}^{(\ell)} = \sum_{\alpha \in \mathcal{P}_{\leq \ell}} \frac{\partial A_i^{(\ell)}}{\partial W_\alpha} \frac{\partial A_j^{(\ell)}}{\partial W_\alpha}$$

and the first-order relation  $\Delta \mathbf{A}^{(\ell)} \approx -\eta \Phi^{(\ell)} \nabla_{\mathbf{A}^{(\ell)}} \mathcal{L}$  still hold for any differentiable sequential feedforward architecture (i.e. one without skip connections). For other modules the same logic applies: there is an own-module contribution from parameters in layer  $\ell$  and a propagated earlier-layer contribution obtained by sandwiching  $\Phi^{(\ell-1)}$  between the appropriate Jacobians. Only the explicit forms of  $D_\ell$  and  $J_\ell$  change.